


Introduction to Parallel Computing

Martin Čuma
Center for High Performance Computing
University of Utah
m.cuma@utah.edu



- Types of parallel computers.
- Parallel programming options.
- OpenMP, OpenACC, MPI
- Higher level languages
- Debugging, profiling and libraries
- Summary, further learning.



- **Faster CPU** clock speed
 - Higher voltage = more heat – not sustainable
- Work distribution
 - **Vectorization** – process more than one value at a time
 - **Parallelization** – spread work over multiple processing elements
 - Specialization – application specific processors (ASIC), programmable logic (FPGA)



Single processor:


- SISD – single instruction single data.

Multiple processors:


- SIMD - single instruction multiple data.
- MIMD – multiple instruction multiple data.
 - Shared Memory
 - Distributed Memory
- Current processors combine SIMD and MIMD
 - Multi-core CPUs w/ SIMD instructions (AVX, SSE)
 - GPUs with many cores and SIMT

- All processors have access to local memory
- Simpler programming
- Concurrent memory access
- More specialized hardware
- Representatives:
Linux clusters nodes 8-32
core GPU nodes


Dual quad-core node



Many-CPU node (e.g. SGI)



- Process has access only to its local memory
- Data between processes must be communicated
- More complex programming
- Cheap commodity hardware
- Representatives:
Linux clusters




8 node cluster (64 cores)

- Process (task)

Entity that executes a program – has its own memory space, execution sequence, is independent from other processes

- Thread

Has own execution sequence but shares memory space with the original process - a process may have many threads



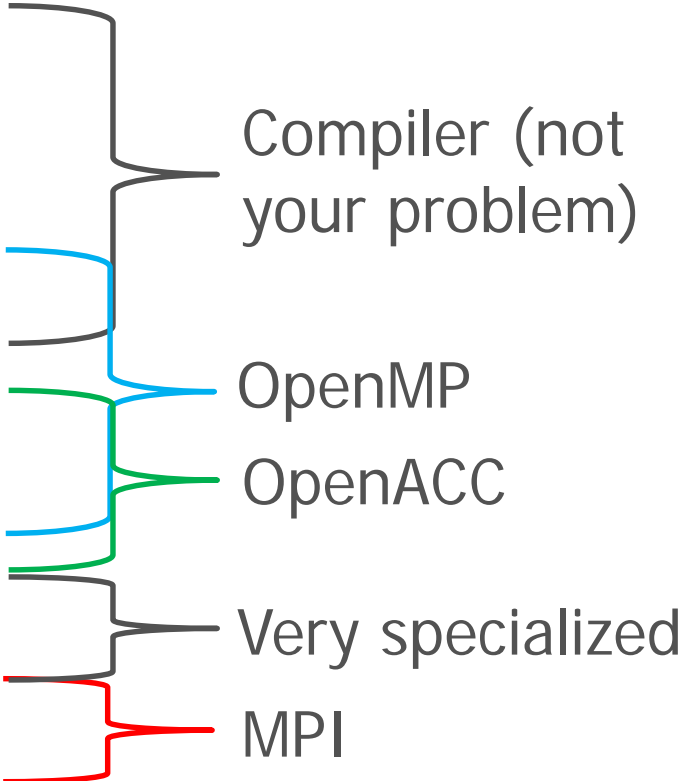
Shared Memory

- Threads
 - POSIX Pthreads, **OpenMP** (CPU, MIC), **OpenACC**, CUDA (GPU)
- Processes
 - message passing, independent processes

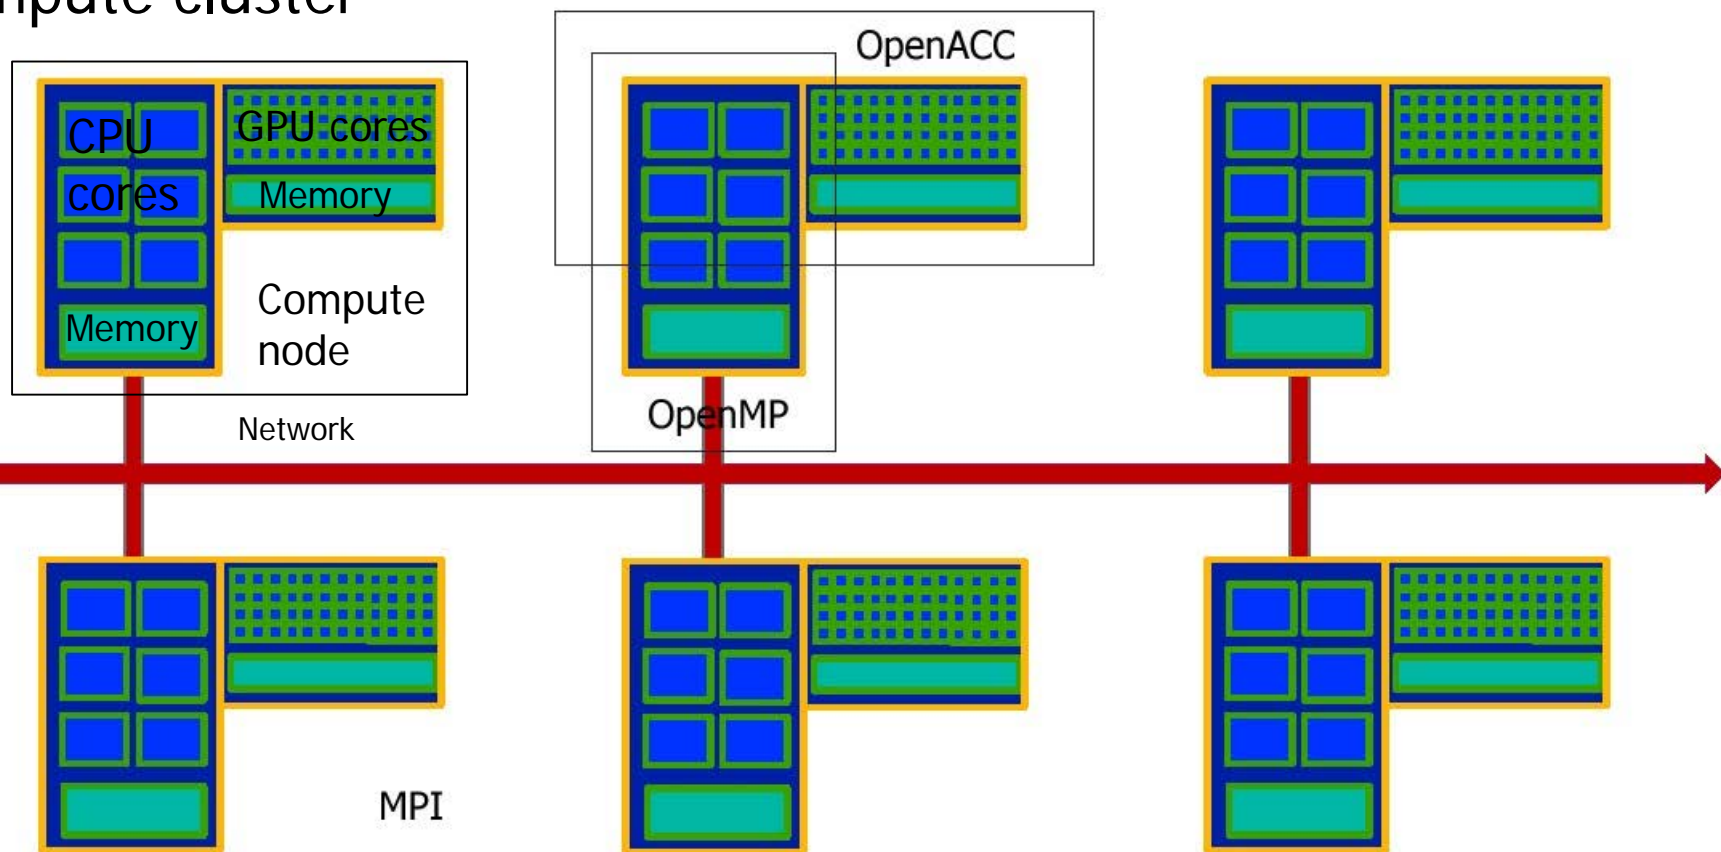
Distributed Memory

- Independent processes
- Message passing libraries
 - General – **MPI**, PVM, language extensions (Co-array Fortran, UPC. ...)

Higher level programming languages (Python, R, Matlab) do a combination of these approaches under the hood.

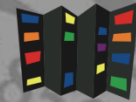
- Instruction level (ILP)
 - Instruction pipelining, speculative execution, branch prediction, ...
 - Vector (SIMD)
 - Multi-core/Multi-socket SMP
 - Accelerators (GPU, MIC)
 - FPGA, ASIC
 - Distributed clusters
- 
- Compiler (not your problem)
- OpenMP
- OpenACC
- Very specialized
- MPI

Compute cluster



- Compiler directives to parallelize (CPU or GPU)
 - Fortran – source code comments
`!$omp parallel/!$omp end parallel`
 - C/C++ - #pragmas
`#pragma omp parallel`
- Small set of subroutines
- Degree of parallelism specification
 - `OMP_NUM_THREADS` or
`omp_set_num_threads(INTEGER n)`

- Compiler directives to offload to GPU
 - Fortran – source code comments
`!$acc kernels/!$acc end kernels`
 - C/C++ - #pragmas
`#pragma acc kernels`
- Small set of subroutines
- Degree of parallelism specification
 - `OMP_NUM_THREADS` or
`omp_set_num_threads(INTEGER n)`



- Communication library
- Language bindings:
 - C/C++ - `int MPI_Init(int argv, char* argc[])`
 - Fortran - `MPI_Init(INTEGER ierr)`
- Quite complex (100+ subroutines)
but only small number used frequently
- User defined parallel distribution

- saxpy – vector addition: $\bar{z} = a\bar{x} + \bar{y}$
- simple loop, no cross-dependence, easy to parallelize

```
subroutine saxpy_serial(z, a, x, y, n)
integer i, n
real z(n), a, x(n), y(n)

do i=1, n
    z(i) = a*x(i) + y(i)
enddo
return
```

```
subroutine saxpy_parallel_omp(z, a, x, y, n)
integer i, n
real z(n), a, x(n), y(n)
```

```
!$omp parallel do
```

```
do i=1, n
```

```
    z(i) = a*x(i) + y(i)
```


```
enddo
```

```
return
```

```
$ gcc -fopenmp saxpy.f
```

```
$ export OMP_NUM_THREADS=16
```

```
$ ./a.out
```





- Data dependencies
 - Private (thread-local) variables
 - Flow dependence – rearrangement
 - Reduction (sum over threads)
- Scheduling
 - What runs on what thread – schedule, task,...
- Advanced features
 - Thread affinity (to CPU core)
 - Vectorization
 - Accelerator offload

```
x = a(i)  
b(i) = c + x
```

```
a(i) = a(i+1) + x
```

```
x += a(i)
```



```
subroutine saxpy_parallel_oacc(z, a, x, y, n)
integer i, n
real z(n), a, x(n), y(n)
```

```
!$acc kernels datain(x,y) dataout(z)
```

```
do i=1, n
```

```
    z(i) = a*x(i) + y(i)
```


```
enddo
```

```
return
```

```
$ pgcc -acc -Minfo=accel saxpy.f
```

```
$ pgacceleinfo To verify that GPU is available
```

```
$ ./a.out
```



- Data dependencies (Like in OpenMP)
- Data locality
 - Transfers from host to GPU and back take time
 - need to minimize them

```
#pragma acc data [copyin, copyout, create,...]
```
- Parallel regions
 - More explicit execution control (warps, threads)

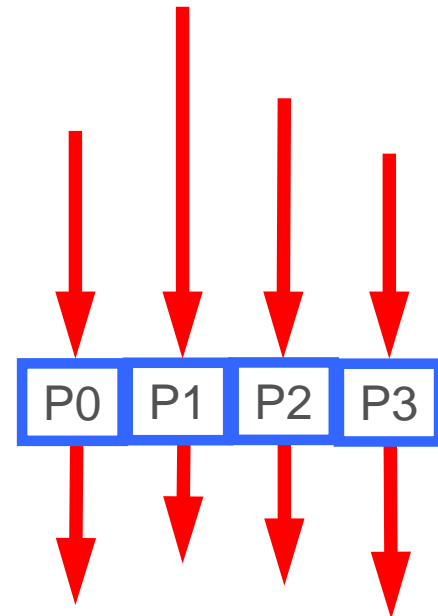
```
#pragma acc parallel
```
- Procedure calls
 - If procedure is executed on the GPU

```
#pragma acc routine
```

```
subroutine saxpy_parallel_mpi(z, a, x, y, n)
integer i, n, ierr, my_rank, nodes, i_st, i_end
real z(n), a, x(n), y(n)
```

```
call MPI_Init(ierr)
call MPI_Comm_rank(MPI_COMM_WORLD, my_rank, ierr)
call MPI_Comm_size(MPI_COMM_WORLD, nodes, ierr)
i_st = n/nodes*my_rank+1
i_end = n/nodes*(my_rank+1)
```

```
do i=i_st, i_end
    z(i) = a*x(i) + y(i)
enddo
call MPI_Finalize(ierr)
return
```

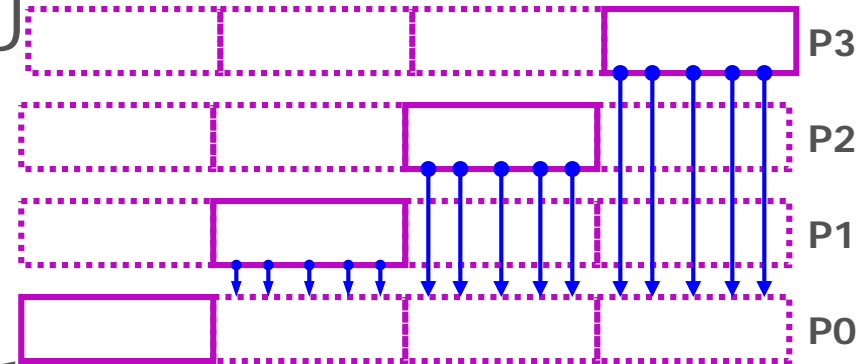


$z(i)$ operation on 4 processes (tasks)

$z(1)$... $n/4$)	$z(n/4+1)$... $2*n/4$)	$z(2*n/4+1)$... $3*n/4$)	$z(3*n/4+1)$... n)
-----------------------	-----------------------------	-------------------------------	---------------------------

- Result on the first CPU

```
include "mpif.h"
integer status(MPI_STATUS_SIZE)
if (my_rank .eq. 0 ) then
  do j = 1, nodes-1
    do i= n/nodes*j+1, n/nodes*(j+1)
      call MPI_Recv(z(i),1,MPI_REAL,j,0,MPI_COMM_WORLD,
&      status,ierr)
    enddo
  enddo
else
  do i=i_st, i_end
    call MPI_Send(z(i),1,MPI_REAL,0,0,MPI_COMM_WORLD,ierr)
  enddo
endif
```



- Collective communication

```
real zi(n)
```

```
j = 1
```

```
do i=i_st, i_end
```

```
    zi(j) = a*x(i) + y(i)
```

```
    j = j + 1
```

```
enddo
```

```
call MPI_Gather(zi,n/nodes,MPI_REAL,z,n/nodes,MPI_REAL,
&
0,MPI_COMM_WORLD,ierr)
```

Send data


Receive data

Root process

- Result on all nodes

```
call MPI_AllGather(zi,n/nodes,MPI_REAL,z,n/nodes,
&
MPI_REAL,MPI_COMM_WORLD,ierr)
```

No root process





- Explicit task based parallelism
 - manual work distribution
 - task communication and synchronization
- Communication patterns
 - due to different data distribution
- Many advanced features
 - blocking vs. non-blocking communication
 - derived data types
 - topologies
 - ...

broadcast
reduction
gather/scatter
...

- Different networks
 - Ethernet
 - InfiniBand
 - Intel OmniPath
 - most MPI distributions now come with multiple networks support
- Several distributions follow the MPI standard
 - MPICH, MVAPICH2
 - Intel MPI, Cray MPI,...
 - OpenMPI
 - Ensure that build and run is done with the same distribution (ABI compatibility)

But wait, my program is not in C or Fortran




Interpreted languages are popular

- Matlab, Python, R

Each has some sort of parallel support, but most likely it will not perform as well as using OpenMP or MPI with C/Fortran.

Try to parallelize (and optimize) your Matlab/Python/R code and if it's still not enough consider rewriting in C++ or Fortran.

- Using parallelization in the program run through interactive or batch job
 - multi-threading and/or multi-processing packages (parfor, mpi4py, R parallel, Rmpi, ...)
- Using built in job submission
 - Matlab Parallel Server, rslurm, python Dask, snakemake
- Independent calculations in parallel
 - launching concurrent calculations in a job



Threads

- Built in Matlab functions. Vector/matrix operations threaded (and vectorized) through Intel MKL library, many other functions also threaded

Tasks (processes)

- *Parallel Computing Toolbox* allows for task based parallelism
- *Parallel Server* can distribute tasks to multiple nodes
- Great for independent calculations, when communication is needed uses MPI under the hood

<https://www.chpc.utah.edu/documentation/software/matlab.php>

- Parallel program

```
function t = parallel_example
parfor idx = 1:16
    A(idx) = idx;
end
```

Will launch loop iterations on multiple workers

- Parallel worker pool on a single machine

```
poolobj=parpool('local',8);
parallel_example;
delete(poolobj);
```

Starts multiple workers pool

- Parallel pool on a cluster

```
c = parcluster;
c.AdditionalProperties.QueueName = 'kingspeak';
...
j = c.batch(@parallel_example, 1, {}, 'Pool', 4);
j.State
j.fetchOutputs{:}
```

Submits cluster job

- Parallel worker pool on a single node
 - best run from a SLURM job
[loop_parallel_onenode.m](#), [run_matlab_onenode.m](#),
[run_matlab_onenode.slr](#)
 - <https://git.io/CHPC-Intro-to-Parallel-Computing-Matlab>
 - `sbatch run_matlab_onenode.slr`
- Parallel worker pool on a multiple nodes
 - must run from inside of Matlab
 - start Matlab on interactive node inside of a FastX session
`ml matlab`
`matlab &`
 - [loop_parallel.m](#), [parallel_multinode.m](#)
`parallel_multinode`

Threads

- Under the hood threading with specially built (or Microsoft) R for vector/matrix operations using MKL
- *parallel* R library

Tasks (processes)


- *parallel* R library (uses *multicore* for shared and *snow* for distributed parallelism)
- Parallelized **apply* functions, e.g. *mclapply*
- *Rmpi* library provides MPI like functionality
- Many people run multiple independent R instances in parallel

<https://www.chpc.utah.edu/documentation/software/r-language.php>

- Load libraries

```
library(parallel)  
library(foreach)  
library(doParallel)
```

hostlist.txt comes from a job script
srun -n \$SLURM_NTASKS > hostlist.txt



- Start R cluster

```
hostlist <- paste(unlist(read.delim(file="hostlist.txt",  
header=F, sep = " ")))  
cl <- makeCluster(hostlist)  
registerDoParallel(cl)  
clusterEvalQ(cl, .libPaths("/uufs/chpc.utah.edu/sys/installdir/  
Rlibs/3.5.2i"))
```

this is only needed if running on multiple nodes


- Run parallel loop

```
r <- foreach(icount(trials), .combine=rbind) %dopar% {}
```

- Stop R cluster

```
stopCluster(cl)
```

- Parallel R on one node
 - best run from a SLURM job
[parallel-onenode-iris.R](#), [R-parallel-onenode-iris.slr](#)
 - <https://git.io/CHPC-Intro-to-Parallel-Computing-R>
 - `sbatch R-parallel-onenode-iris.slr`
- Parallel R multiple nodes
 - must specify list of nodes where R workers run
[parallel-multinode-iris.R](#), [R-parallel-multinode-iris.slr](#)
 - `sbatch R-parallel-onenode-iris.slr`
- Submit SLURM job directly from R - `rslurm`
 - SLURM-aware apply function, some issues with results collection
 - [rslurm-example.R](#)



Threads

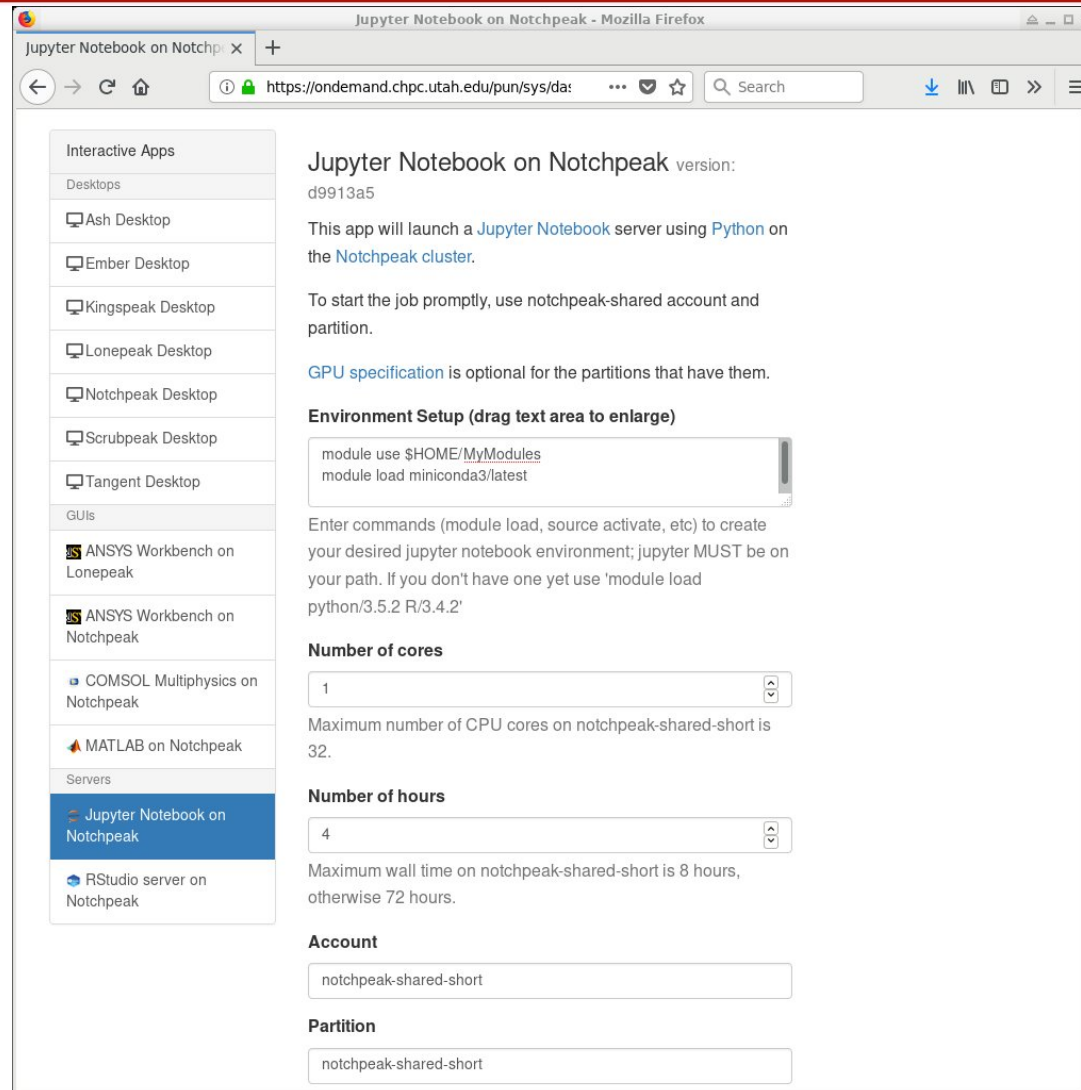
- No threads in Python code because of GIL (Global Interpreter Lock)
- C/Fortran functions can be threaded (e.g. *NumPy* - Anaconda)

Tasks (processes)

- Several libraries that use MPI under the hood, most popular is *mpi4py*
- More-less MPI function compatibility, but slower communication because of the extra overhead
- Also many other data-parallel libraries, e.g. *Dask*

<https://www.chpc.utah.edu/documentation/software/python.php>

- Several options listed at <https://www.chpc.utah.edu/documentation/software/jupyterhub.php>
- The easiest is to use Open OnDemand



Jupyter Notebook on Notchpeak - Mozilla Firefox

Jupyter Notebook on Notchpeak x +

https://ondemand.chpc.utah.edu/pun/sys/da: Search

Interactive Apps

Desktops

- Ash Desktop
- Ember Desktop
- Kingspeak Desktop
- Lonepeak Desktop
- Notchpeak Desktop
- Scrubpeak Desktop
- Tangent Desktop

GUIs

- ANSYS Workbench on Lonepeak
- ANSYS Workbench on Notchpeak
- COMSOL Multiphysics on Notchpeak
- MATLAB on Notchpeak

Servers

- Jupyter Notebook on Notchpeak
- RStudio server on Notchpeak

Jupyter Notebook on Notchpeak version: d9913a5

This app will launch a Jupyter Notebook server using Python on the Notchpeak cluster.

To start the job promptly, use notchpeak-shared account and partition.

GPU specification is optional for the partitions that have them.

Environment Setup (drag text area to enlarge)

```
module use $HOME/MyModules
module load miniconda3/latest
```

Enter commands (module load, source activate, etc) to create your desired jupyter notebook environment; jupyter MUST be on your path. If you don't have one yet use 'module load python/3.5.2 R/3.4.2'

Number of cores

1

Maximum number of CPU cores on notchpeak-shared-short is 32.

Number of hours

4

Maximum wall time on notchpeak-shared-short is 8 hours, otherwise 72 hours.

Account

notchpeak-shared-short

Partition

notchpeak-shared-short

- Our personal favorite is to ignore all the Python parallel efforts, divide the data into independent parts and run multiple Python processes on parts of the data concurrently
- Only works if data can be split
- Use various approaches for independent parallel calculations listed at <https://www.chpc.utah.edu/documentation/software/serial-jobs.php>
- More on this later

- With relatively small effort one can use Dask
- Install Miniconda

```
wget https://repo.continuum.io/miniconda/Miniconda3-latest-Linux-x86_64.sh
```

```
bash ./Miniconda3-latest-Linux-x86_64.sh -b -p  
$HOME/software/pkg/miniconda3
```

```
mkdir -p $HOME/MyModules/miniconda3
```

```
cp  
/uufs/chpc.utah.edu/sys/installdir/python/modules/miniconda3/latest.lua  
$HOME/MyModules/miniconda3
```

- Use own miniconda and install Jupyter and Dask

```
module use $HOME/MyModules  
module load miniconda3/latest  
conda install jupyter dask
```

- Start Open OnDemand Jupyter notebook
 - log into ondemand.chpc.utah.edu with CHPC credentials



- Go to Interactive Apps - Jupyter Notebook on notchpeak
- In the Environment Setup text box, put:

```
module use $HOME/MyModules  
module load miniconda3/latest
```
- Use notchpeak-shared-short for account and partition, and select your choice of CPU cores and walltime hours (within the listed limits). Then hit Launch to submit the job.
- Once the job starts, hit the blue Connect to Jupyter button
- Open one of the following notebooks:
[dask_embarrass.ipynb](#), [dask_slurmcluster.ipynb](#),
[dask_slurm_xarray.ipynb](#)
- DASK also allows to submit jobs to SLURM (last 2 examples)



Independent calculations

- Different approaches based on the nature of the calculations
 - Runtime length, variability, number of calculations
- Similar runtime, small calculation count
 - Shell script in a SLURM job

```
#!/bin/bash
for (( i=0; i < $SLURM_NTASKS ; i++ )); do
    /path_to/myprogram $i &
done
wait
```
 - `srun --multi-prog`

```
srun --multi-prog my.conf
cat my.conf
0-11 ./example.sh %t
```

<https://www.chpc.utah.edu/documentation/software/serial-jobs.php>

- Mini-scheduler inside of a job
 - to launch calculations till all are done
 - GNU Parallel - <https://www.gnu.org/software/parallel/>
 - TACC Launcher - <https://www.tacc.utexas.edu/research-development/tacc-software/the-launcher>
 - CHPC Submit - <https://www.chpc.utah.edu/documentation/software/serial-jobs.php#submit>
- Workflow managers
 - More on this later
 - Makeflow, Swift, Snakemake, Pegasus
- Distributed computing resources
 - Open Science Grid - <https://opensciencegrid.org/>



- Useful for finding bugs in programs
- Several free
 - `gdb` – GNU, text based, limited parallel
 - `ddd` – graphical frontend for `gdb`
- Commercial that come with compilers
 - `pgdbg` – PGI, graphical, parallel but not intuitive
 - `pathdb`, `idb` – Pathscale, Intel, text based
- Specialized commercial
 - `totalview` – graphical, parallel, CHPC has a license
 - `ddt` - Distributed Debugging Tool
 - **Intel Inspector XE** – memory and threading error checker
- How to use:
 - http://www.chpc.utah.edu/docs/manuals/software/par_devel.html

- Parallel debugging more complex due to interaction between processes
- DDT is the debugger of choice at CHPC
 - Expensive but academia get discount
 - How to run it:
 - compile with `-g` flag
 - run `ddt` command
 - fill in information about executable, parallelism, ...

- Details:

<https://www.chpc.utah.edu/documentation/software/debugging.php>

- Further information

<https://www.allinea.com/products/ddt>

The screenshot displays the Allinea DDT debugger interface. The main window shows a C program named `watchmatrix.c` with the following code:

```

35  for (i = 0; i < 30; i++)
36      for (j = 0; j < 30; j++)
37          A[i][j] = 0;
38  for (i = 0; i < 30; i++)
39      for (j = 0; j < 30; j++)
40          B[i][j] = 0;
41  for (i = 0; i < 30; i++)
42      for (j = 0; j < 30; j++)
43          C[i][j] = 0;
44  for (i = 0; i < 30; i++)
45      for (j = 0; j < 30; j++)
46          C[i][j] = A[i][j] + B[i][j];
47
48


```

The `watchmatrix.c` file is open in the editor, and the `watchmatrix.c` file is selected in the Project Files pane. The `watchmatrix.c` file is selected in the Project Files pane.


A `DDT - Edit VIspoint` dialog box is open, showing the location of the watchpoint. The `Location` tab is selected, and the `File` is `curial/doc/training/programs/watchpoint/watchmatrix.c`. The `Line Number` is `41`. The `Visualise` tab is also visible, showing the `Mesh Type` as `Rectilinear` and the `Variable Centering` as `Zone`. The `Array Expression` is `C[i][j]`.


A `Window 1` visualization window is open, displaying a 2D heatmap of the data. The title bar indicates `DB: 001324225388.ftables-ddt.sim2` and `Cycle: 268 Time: 268`. The plot shows a color gradient from blue (low values) to red (high values). The x-axis is labeled `i` and the y-axis is labeled `j`. The plot area is divided into four quadrants labeled A, B, C, and D. The color scale on the left ranges from `Min: 0.000` to `Max: 0.010`.

The `Visualisation Points` pane at the bottom shows the current state of the visualization. The `Processes` pane shows `All` and the `Threads` pane shows `all`. The `File` pane shows `watchmatrix.c` and the `Line` pane shows `41`.



- Measure performance of the code
- Serial profiling
 - discover inefficient programming
 - computer architecture slowdowns
 - compiler optimizations evaluation
 - gprof, pgprof, pathopt2, Intel tools
- Parallel profiling
 - target is inefficient communication
 - **Intel Trace Collector and Analyzer, AdvisorXE, VTune**



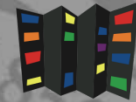


- Use libraries for common operationsSerial
 - BLAS, LAPACK – linear algebra routines
 - MKL, ACML – hardware vendor libraries
- Parallel
 - ScaLAPACK, PETSc, FFTW
 - MKL – dense and sparse matrices
- Design a new code around existing library
 - PETSc, Trilinos,...

- MPICH, MVAPICH2 and Intel MPI are cross-compatible using the same ABI
 - Can e.g. compile with MPICH on a desktop, and then run on the cluster using MVAPICH2 and InfiniBand
- Intel and PGI compilers allow to build "unified binary" with optimizations for different CPU platforms
 - But in reality it only works well under Intel compilers
- On a desktop

```
module load intel mpich
mpicc -xCORE-AVX512,CORE-AVX2,AVX program.c -o program.exe
mpirun -np 4 ./program.exe
```
- On a cluster

```
srun -N 2 -n 24 ...
module load intel mvapich2
mpirun -np $SLURM_NTASKS ./program.exe
```
- <https://www.chpc.utah.edu/documentation/software/single-executable.php>



- Shared vs. Distributed memory parallelism
- OpenMP, OpenACC and MPI for low level parallelism
- Different approaches for higher level languages
- Many ways to run independent calculations in parallel
- There are tools for debugging, profiling

- CHPC lectures
 - <https://www.chpc.utah.edu/presentations/index.php>
- XSEDE HPC Summer Boot Camp
 - OpenMP, OpenACC, MPI
 - <https://www.youtube.com/XSEDETraining>
- Petascale Computing Institute
 - Wide range of parallel programming topics
 - videos at <https://bluewaters.ncsa.illinois.edu/bw-petascale-computing-2019/agenda>
- XSEDE online training
 - <https://www.xsede.org/web/xup/online-training>